

Programming Z3

Nikolaj Bjørner
Microsoft Research

SMT workshop
July 8, 2019

<https://z3examples-nbjorner.notebooks.azure.com/j/notebooks>

<https://tinyurl.com/y3r67rd6>

SMTWorkshop.ipynb

Outline

- Past, present
 - An update on Z3
 - Some applications
- Some not so Secret Sauce
- Actually Programming Z3
- Active Directions

An update



The Z3 Theorem Prover

Manage topics

10,830 commits | 8 branches | 14 releases | 121 contributors | View license

Branch: master | New pull request | Create new file | Upload files | Find file | Clone or download

Nikolaj Bjørner add #2298 to regression/example ... Latest commit 25:9341 31 minutes ago

File	Commit Message	Time Ago
cmake	Change from BINARY_DIR to PROJECT_BINARY_DIR	14 days ago
contrib	Fix bug in qprofdiff	4 months ago
doc	Change from BINARY_DIR to PROJECT_BINARY_DIR	14 days ago
examples	add #2298 to regression/example	31 minutes ago
noarch	follow instructions from #1879	8 months ago
	Updated nuget package spec and directions	7 months ago
	Fix z3 static link options	6 days ago
	add #2298 to regression/example	31 minutes ago
	[TravisCI] Implement TravisCI build and testing infrastructure for Linux	2 years ago
	set text default to auto to try to avoid crlf disasters	5 years ago
	incrementally adding files from dotnet core pull request from @yattli	4 months ago
	Revert "api: dotnet: switch to multi-targeting project and modern cma...	4 months ago
	Change from BINARY_DIR to PROJECT_BINARY_DIR	14 days ago
	update license for space/quotes per #982	2 years ago
CMake.md	merge with Z3Prover/master	11 months ago
md	add mace3 status	6 months ago
NOTES	release notes	5 months ago
	merge with Z3Prover/master	11 months ago
	merge with Z3Prover/master	11 months ago



a theorem prover from Microsoft Research. It is licensed under the MIT license.

are not familiar with Z3, you can start [here](#).

Build binaries for releases are available from [here](#), and nightly builds from [here](#).

can be built using [Visual Studio](#), a [Makefile](#) or using [CMake](#). It provides bindings for several programming languages.

[the release notes](#) for notes on various stable releases of Z3

[Build status](#)

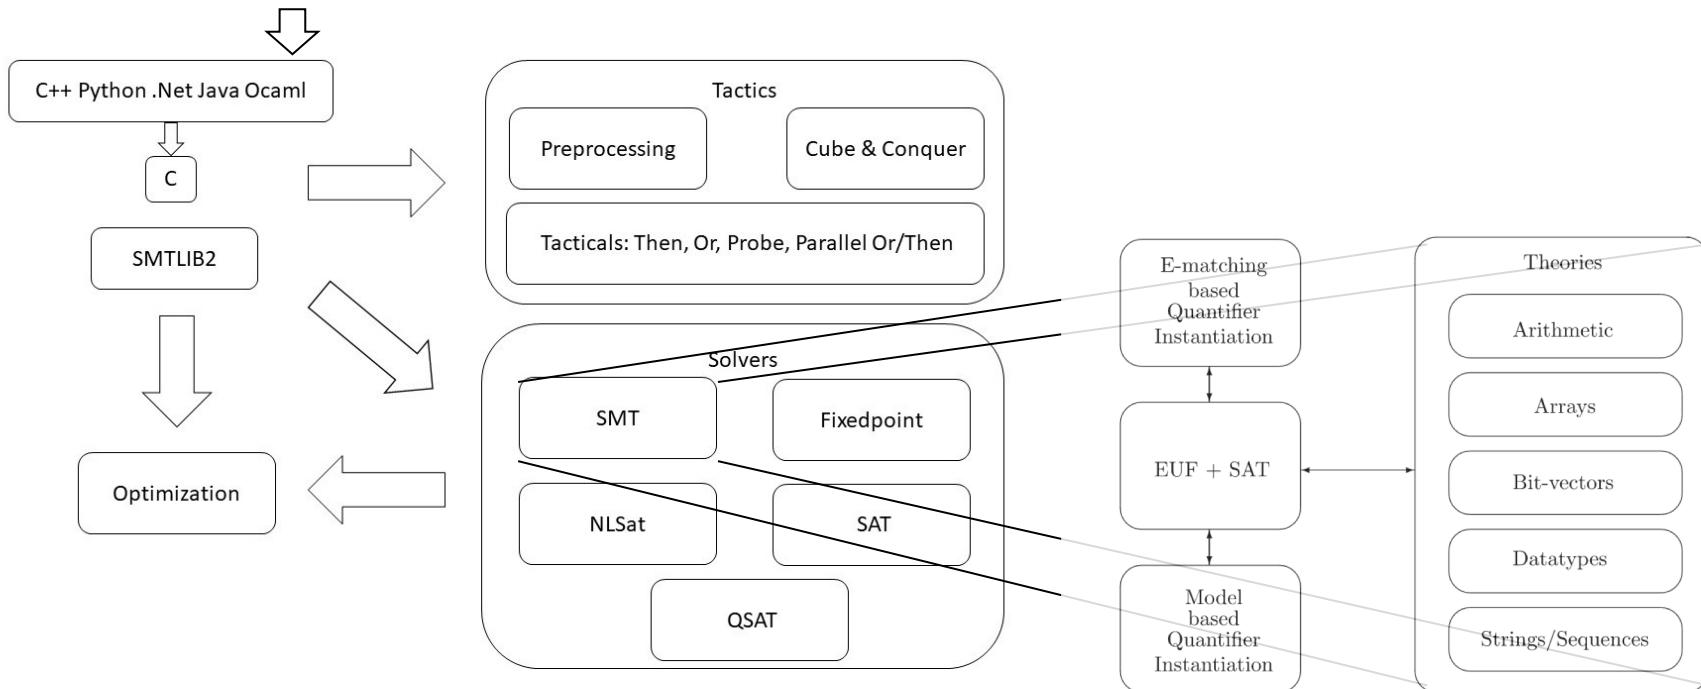
Windows x64	Windows x86	Windows x64	Ubuntu x64	Debian x64	macOS	TravisCI



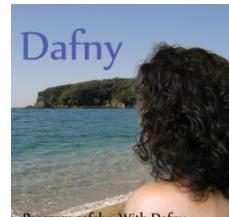
<https://github.com/Z3Prover/z3>

Integrates a wealth of domains and solvers

NuGet, pip, Opam,
Rust, JavaScript, Haskell



Symbolic Analysis Engines



TERMINATOR

SLAYER
HAVOC

BOOGIE

Relevancy Propagation

Efficient E-matching for SMT solvers

Z3 Internals

SLS, floats

νZ : Opt+MaxSMT

μZ : Datalog

Generalized PDR

Existential Reals

Model Constructing SAT

CutSAT: Linear Integer Formulas

Quantified Bit-Vectors

Linear Quantifier Elimination

Model Based Quantifier Instantiation

Generalized, Efficient Array Decision Procedures

Engineering DPLL(T) + Saturation

Effectively Propositional Logic

Model-based Theory Combination.

Some Microsoft Uses of **Z3**

also: Dynamics Tax tool, Visual Studio C++ compiler, Azure Blockchain, Static Driver Verifier, Pex

Microsoft Security Risk Detection

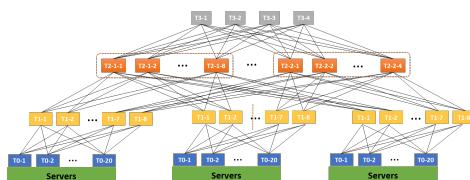


Solved billions of fuzzing constraints

Program paths

Bit-vector + array logic

SecGuru and FIB verifier



Online checks of $O(10^5)$ Azure routers + ACLs

Network Configurations

Bit-vector logic

Dynamics Product Configurator



Maintains design space of parameters

Production Line Configs

Enumerate Consequences

Project Everest



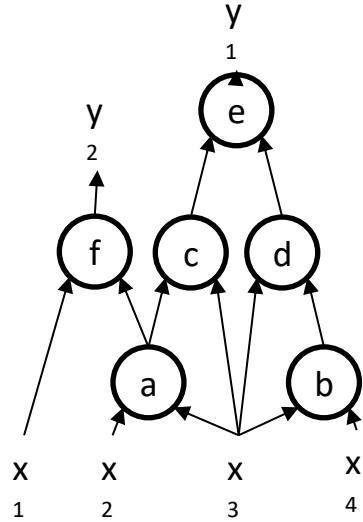
Formal proofs for WinQUIC

Program + Spec

$\forall x \exists y p(x) \rightarrow x \geq y$

Quantum: Reversible pebbling game

Example: find a pebbling strategy using 6 pebbles.



pebbling configurations

$$P_1 = \{\emptyset\},$$

$$P_2 = \{a\},$$

$$P_3 = \{a, b\},$$

$$P_4 = \{a, b, c\},$$

$$P_5 = \{a, b, c, d\},$$

$$P_6 = \{a, b, c, d, e\},$$

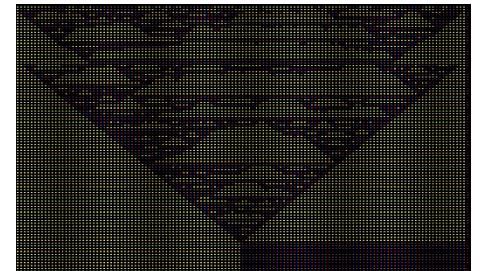
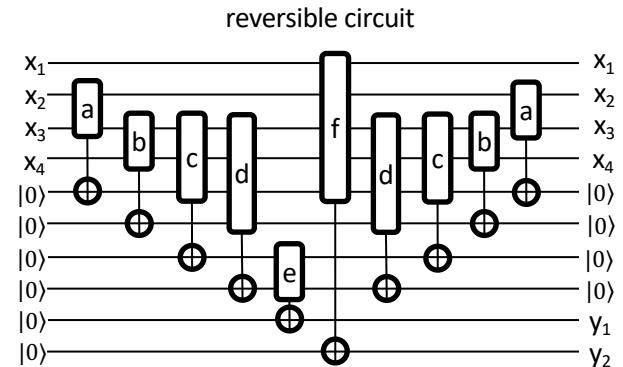
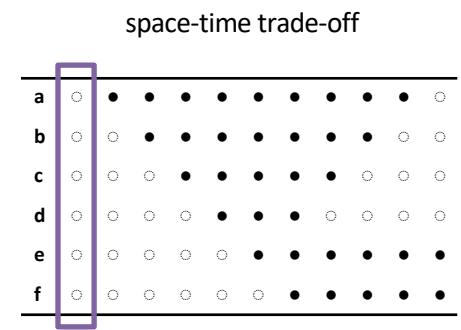
$$P_7 = \{a, b, c, d, e, f\},$$

$$P_8 = \{a, b, c, e, f\},$$

$$P_9 = \{a, b, e, f\},$$

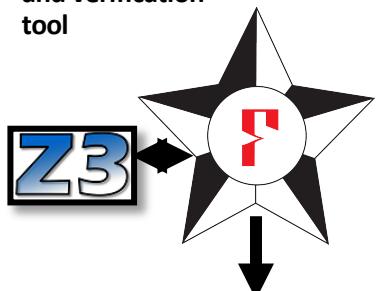
$$P_{10} = \{a, e, f\},$$

$$P_m = P_{11} = \{e, f\}$$



Everest, EverCrypt, EverParse

F*: A general purpose
programming
language
and verification
tool



Compiler from
(a subset of)
F* to C

C

Math spec in F*

`poly1305_mac` computes a polynomial in $GF(2^{130}-5)$, storing the result in `tag`, and not modifying anything else

```
val poly1305_mac: tag:nbytes 16 →  
    len:u32 →  
    msg:nbytes len{disjoint tag msg} →  
    key:nbytes 32 {disjoint msg key ∧ disjoint tag  
    →  
    ST unit  
    (requires (λ h → msg ∈ h ∧ key ∈ h ∧ tag ∈ h))  
    (ensures (λ h0 _ h1 →  
        let r=Spec.clamp h0.[sub key 0 16] in  
        let s=h0.[sub key 16 16] in  
        modifies {tag} h0 h1 ∧  
        h1.[tag] == Spec.mac_1305 (encode_bytes h0.[msg]) r s))
```

Efficient C implementation

Verification imposes no runtime performance overhead

```
void  
poly1305_mac(uint8_t *tag, uint32_t len, uint8_t *msg, uint8_t *key)  
{  
    uint64_t tmp [10] = { 0 };  
    uint64_t *acc = tmp;  
    uint64_t *r = tmp + (uint32_t)5;  
    uint8_t s[16] = { 0 };  
    Crypto_Symmetric_Poly1305_poly1305_init(r, s, key);  
    Crypto_Symmetric_Poly1305_poly1305_process(msg, len, acc, r);  
    Crypto_Symmetric_Poly1305_poly1305_finish(tag, acc, s);  
}
```

EverCrypt Clients: Mozilla Firefox; WireGuard VPN; Linux Kernel Zinc crypto library; MirageOS unikernel; Tezos blockchain; Microsoft QUIC.

Trusted Financial Software

● I ● | M A N D R A

Imandra is a cloud-native automated reasoning engine.

Imandra's groundbreaking AI helps ensure the algorithms we rely on are safe, explainable and fair.

- Recursive Function Unfolding
- Algebraic ML Datatypes
- Ground Arithmetic

TRY IMANDRA ONLINE

INSTALL IMANDRA LOCALLY

Verifying ReasonReact component logic
— ReasonML & Imandra

4 September 2018

Spotlight on an Imandra user:

In 2017 Aesthetic Integration partnered with Goldman Sachs to help deliver the SIGMA X MTF Auction Book, a

<https://try.imandra.ai/>

Axiomatic Economics

Models of economics formulated using Non-linear Real Arithmetic

Figure 4a. The Laffer curve for transfers

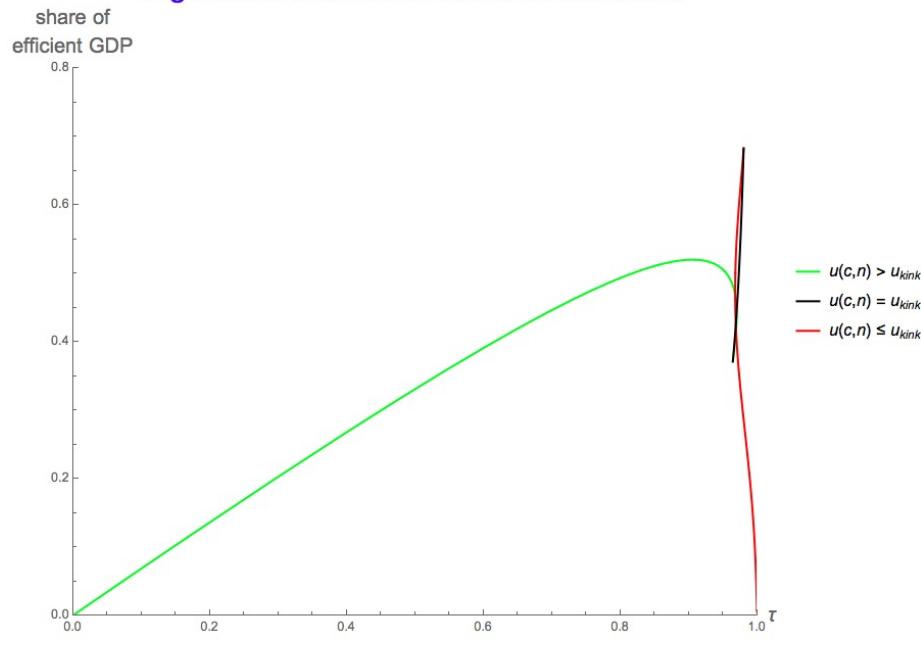
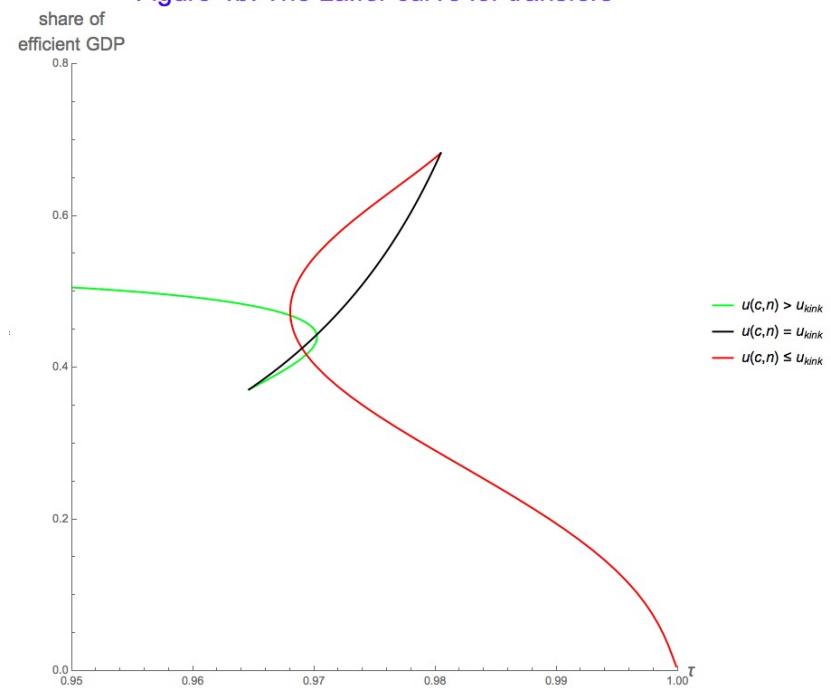


Figure 4b. The Laffer curve for transfers



Casey Mulligan, University of Chicago, School of Economics
Uses Mathematica, Redlog, Z3

Verifying 100Ks Routers in Azure

Intent
Data center layout

Network
Graph
Service

Azure Blob
Storage

Generate
Local Contracts
using NGS

Kusto



Routing Table
Validator



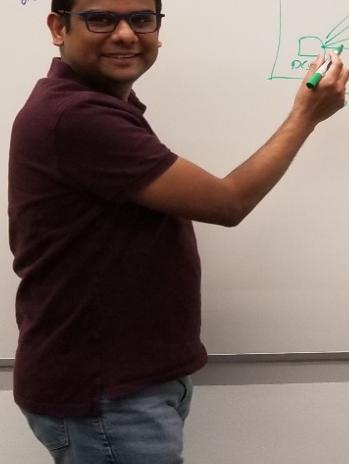
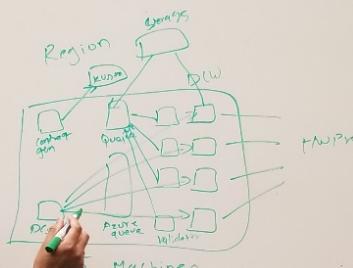
Azure Blob
Storage

Pull forwarding
tables

Hardware
Proxy

Reality
What is on
the routers

- Route Rule
- Prefix
- Next-hop
- Type - static, int



Azure
Queue

Schedule
FIB retrieval

Verification

BlockChain

- Solidity
 - Static Analysis, Leonardo Alt et al
 - Augur - Symbolic Execution
 - SPACER
- runtimeverification.com [Rosu + 30]
 - Verify byte code, ground truth
 - Arithmetic over 256-bit bit-vectors (use linear arithmetic)
- synthetic-minds.com [Srivastava]
 - Recursive functions (modeling a heap library)
- www.certora.com [Grossmann, Katz, Sagiv, Taube]
 - Quantifier Free Arrays + Bit-Vectors (QF_ABV)
 - Quantifier Free Functions + linear arithmetic
 - EPR + linear arithmetic



Rigorous Methods for
Smart Contracts
Dagstuhl, June 1-5 2020
Org: Christakis, B, Maffei, Rosu

What SMT features are used by applications?

- QF_ABV: symex with heap
- QF_S: XML configurations, policies
- UFLIA: Boogie, Everest, Viper, ..
- QF_UFLIA, QF_ABV: Smart contracts
- ALL (Kitchen sink): Pex, Haskell
- Boolean Theories: Operations research applications, approximate counting (Xor)

When can policies **not** be sensitive to strings?

Formulas with 256 bit arithmetic

SOME (NOT SO) SECRET SAUCE

Guiding inferences using models

Model-based Theory Combination [M,B 07]

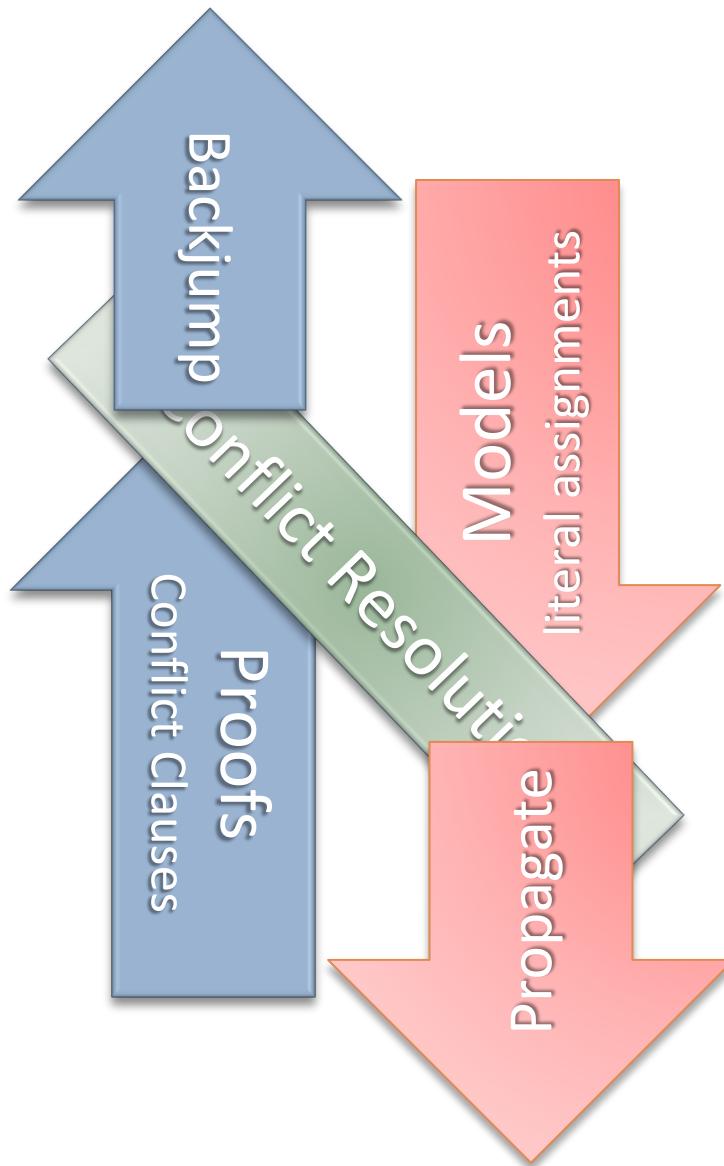
Model-based Quantifier Instantiation [G,M (B) 09]

Generalized Property Directed Reachability [H,B 12]

Model-constructing Satisfiability [J,M 11,12,13]

Model-based Quantifier Elimination [G 14]

Mile High: Modern SAT/SMT search



On Constructing Models

- Build a partial interpretation M by setting an unconstrained variable
- Propagate with M
- Backtrack if propagation leads to a conflict



Role models

On Constructing Models - variant

- Build a partial interpretation M by setting an unconstrained variable
- Extend M by solving sub-formula
- Propagate with M globally
- Backtrack (locally or globally) if propagation leads to a conflict

Model-based Theory Combination

Use a candidate model M_i for a theory T_i and propagate all equalities implied by M_i .

- If $M_i \models T_i \cup \Gamma_i \cup \{ u = v \}$ then propagate $u = v$

Hedging that other theories agree.

Backtrack if some other theory disagrees with $u = v$.

It is cheaper to enumerate equalities for a specific model.

Trick: instrument solver to create as few equalities as possible.

Model-based Quantifier Instantiation

Assume we are given $\psi \wedge \forall x \varphi[x]$,
then use model for ψ as starting point
for search of instantiations of $\forall x \varphi[x]$

```
s.add( $\psi$ )
while True:
    if unsat == s.check():
        return unsat
    M = s.model()
    checker = Solver()
    checker.add( $\neg\varphi^M[x]$ )
    if unsat == checker.check():
        return sat
    M = checker.model()
    find t, such that  $x \notin t, t^M = x^M$ .
    s.add( $\varphi[t]$ )
```

$t^M = x^M$ is not a strict requirement.

It is sufficient to use M to mine for a term t that still satisfies $\varphi[t]$

Model-based Quantifier Elimination

```
def qe(∃vec . F):
    e, a = Solver(), Solver()
    e.add(F)
    a.add(¬F)
    G = False
    while sat == e.check():
        M0 = e.model()
        M1 = [ lit for lit in literals(F) if is_true(M0.eval(lit)) ]
        # assume F is in negation normal form
        assert unsat == a.check(M1)
        M2 = a.unsat_core()
        π = project(M2, vec)
        G = G ∨ π
        e.add(¬π)
    return G
```

project can use M_0 to identify a finite set of solutions to ν that cover all of e

Example: project x from $(y_1 \leq x \dots y_{10} \leq x \wedge x \leq z_1, \dots, z_{10}) =$
 $y_1, \dots, y_9 \leq y_{10}, y_{10} \leq z_1, \dots, z_{10}$

assuming $M(y_1), \dots, M(y_9) \leq M(y_{10})$, corresponds to instantiating x by y_{10}

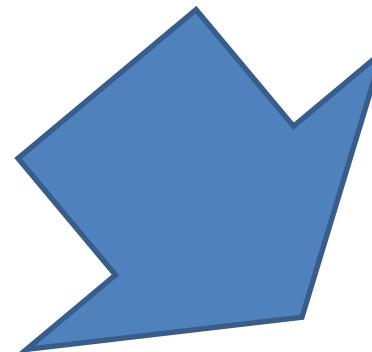
Symbolic model checking as Satisfiability of Horn Clauses

$\text{mc}(x) = x - 10$

if $x > 100$

$\text{mc}(x) = \text{mc}(\text{mc}(x + 11))$ if $x \leq 100$

assert ($x \leq 101 \Rightarrow \text{mc}(x) = 91$)



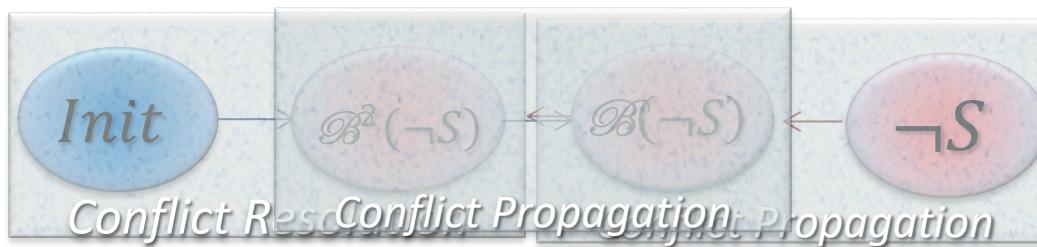
$\forall X. X > 100 \rightarrow \text{mc}(X, X - 10)$

$\forall X, Y, R. X \leq 100 \wedge \text{mc}(X + 11, Y) \wedge \text{mc}(Y, R) \rightarrow \text{mc}(X, R)$

$\forall X, R. \text{mc}(X, R) \wedge X \leq 101 \rightarrow R = 91$

Finds solution for mc

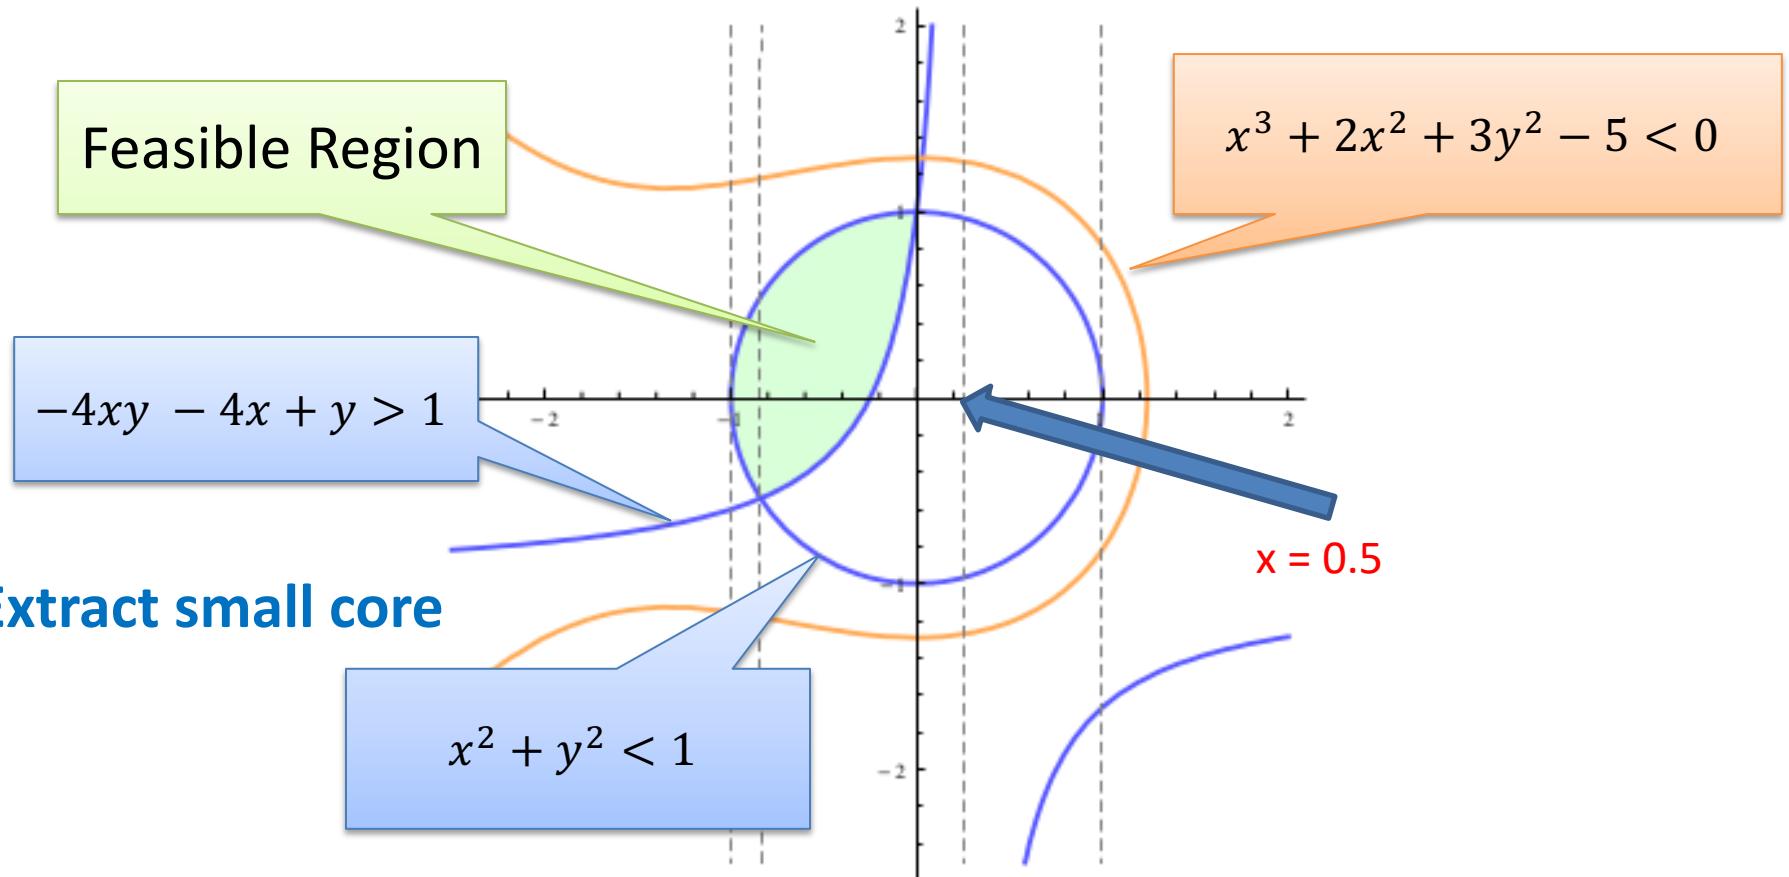
Reachability: Mile-high perspective



- Over-approximate reachable states from Init.
 - Model of Invariant “from above”
- Under-approximate states that can reach $\neg S$
 - Model of trace “from below”
- If over and under-approximations are separated at bound k , produce certificate for separation
 - Strengthens over-approximation of reachable states
 - [SPACER: + under-approximations reachable states]

Solving $\exists R$ Efficiently

A key idea: Use partial solution to guide the search



Strategies

Model-based methods:

- used with back-jumping
- guide inferences

Contrast with

Strategies:

- prune state space of choices
- Limit required inferences



A Basic Quantifier Strategy

```
def strategy(M,j): return  $\bigwedge_{M \neq \text{null}, a \in Atoms, level(j,a) < j} sign(M, a)$ 
def tailv(j): return  $x_{j-1}, x_j, x_{j+1}, \dots$ 

j = 1
M = null
while True:
    if  $F_j \wedge \text{strategy}(M, j)$  is unsat:
        if j == 1:
            return F is unsat
        if j == 2:
            return F is sat
        C = Core( $F_j, \text{strategy}(M, j)$ )
        J = Mbp(tailv(j), C)
        j = index of max variable in J ∪ {1, 2} of same parity as j
         $F_j = F_j \wedge \neg J$ 
        M = null
    else:
        M = current model
        j = j + 1
```

[better ones for QBF by Rabe; Janota, Klieber, B]

Branching Strategies

Can we *train* a branching strategy from benchmarks?

NeuroCore [Daniel Selsam, B. SAT 2019]:

- Train DNN about variables likely to be in core.
- *Periodic refocus* to recalibrate variable splitting queue.
- Claim: DNN provides good precision/recall. Speedups.
- Caveat: DNNs are likely not essential for *refocusing*

Solving formulas

Models, cores

Optimization, Fixed-points

PROGRAMMING Z3 - REALLY

<https://z3examples-nbjorner.notebooks.azure.com/j/notebooks>

<https://tinyurl.com/y3r67rd6>

SOME ACTIVE DIRECTIONS

SM(Boolean Algebras)

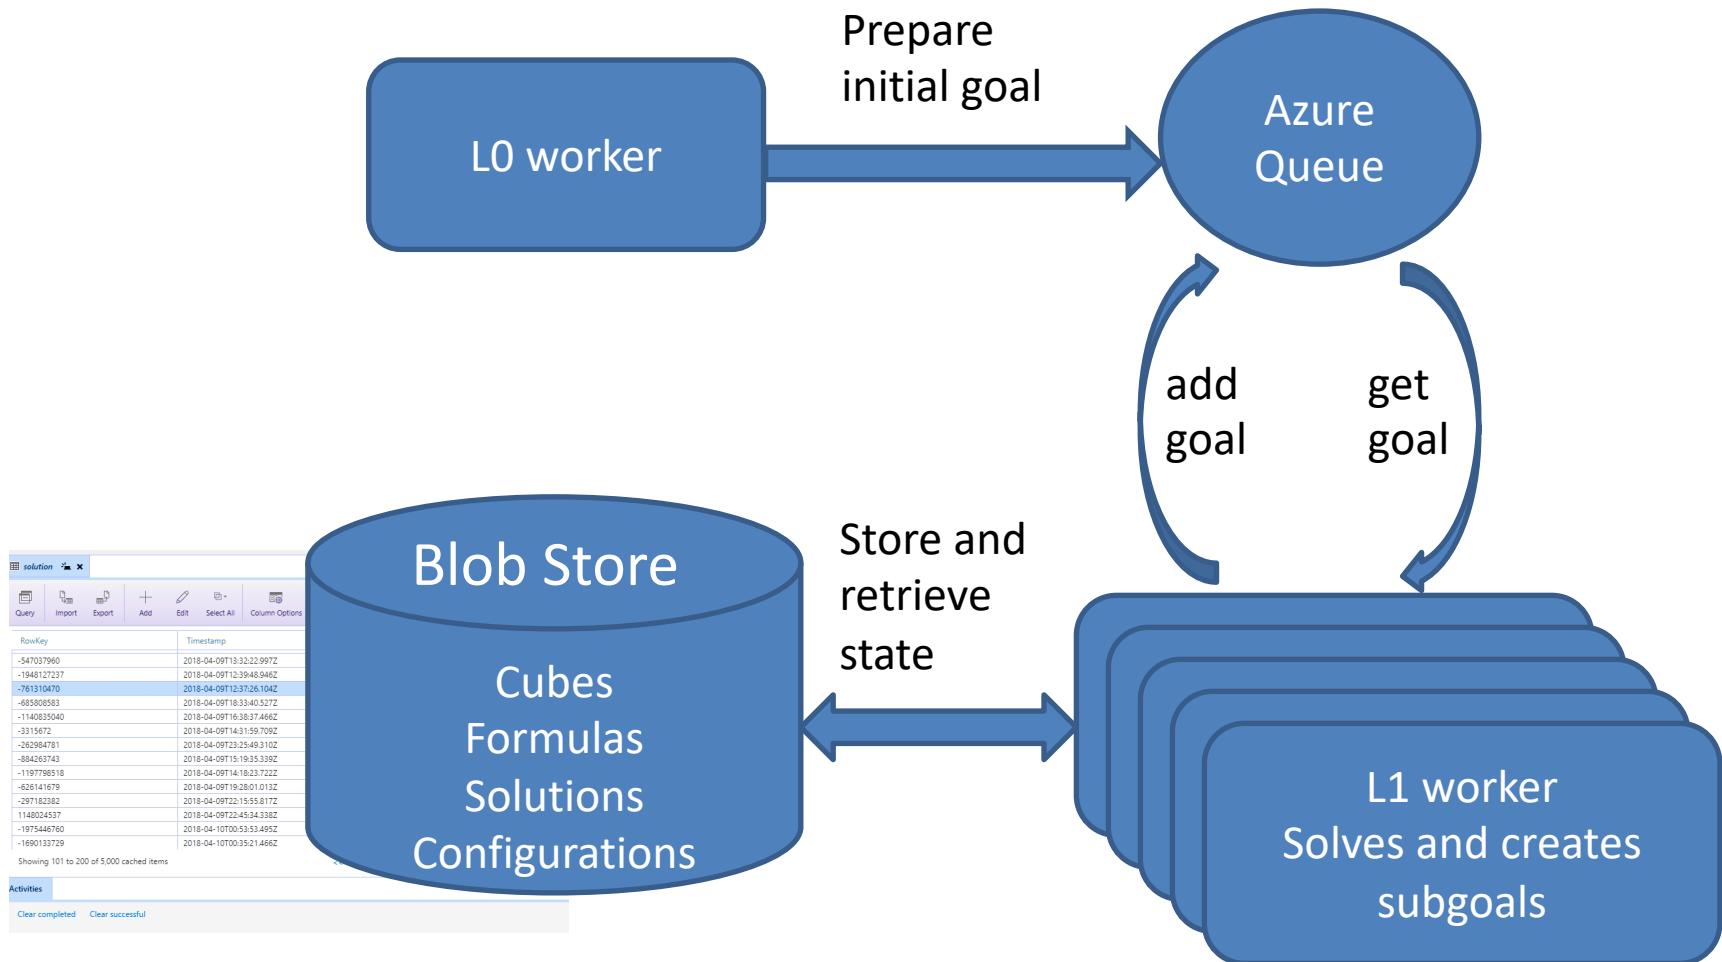
- Plugin in SAT core for
 - Cardinality constraints
 - Pseudo Booleans
 - Xor
- SMT over other Boolean Functions?
 - propagators
 - In-processing
 - conflict analysis

Other recent/in-progress Solvers

- Monoids (strings) and Sequences
- Special Relations (partial, linear orders)
- Transitive Closure as a combinator
- Boolean Algebra and Presburger Arithmetic
- A theory for Job Scheduling
- New full arithmetic solver replacement by Lev Nachmanson

The Cube, the Cloud and Z3

Rahul Kumar (MSR)
Miguel Neves (U Lisboa)



Summary

- SMT applications and the Gartner Hype curve
- Models and Strategies in SMT search
- Programming Z3: BMC, MaxSAT, AllSAT, Cubes
- Directions: Scaling SAT/SMT, CP theories

Thanks: Arie Gurfinkel, Marijn Heule, Rahul Kumar, Leonardo de Moura, Lev Nachmanson, Nina Narodytska, Miguel Angelo Da Terra Neves, Daniel Selsam, and Christoph Wintersteiger.